

# Niladri S. Chatterji

CONTACT INFORMATION	Department of Computer Science Stanford University	<a href="https://niladri-chatterji.github.io/">https://niladri-chatterji.github.io/</a> niladri@cs.stanford.edu
RESEARCH INTERESTS	<p>I am interested in understanding why current deep learning models perform well, with an eye on improving them further.</p> <p><b>Understanding and Building Scaling Laws.</b> Large Language models are trained on internet-scale data and need to tradeoff between the model size and the number of samples under a compute budget. I build tools to help navigate these tradeoffs.</p> <p><b>Robust Deep Learning.</b> Deep learning models are sensitive to shifts in the distribution. I develop techniques to build robust models.</p> <p><b>Generalization and Optimization of Overparameterized Models.</b> Many neural network models generalize well despite perfectly fitting noisy training data. I am interested in understanding what leads to successful generalization and optimization.</p>	
CURRENT POSITION	<b>Postdoctoral Scholar, Stanford University</b> Advisors: Percy Liang & Tatsunori Hashimoto	(June 2021 – Present)
EDUCATION	<b>University of California Berkeley</b> Ph.D. in Physics Advisor: Peter Bartlett and Michael Jordan Thesis: Why do gradient methods work in optimization and sampling? <b>Indian Institute of Technology Bombay</b> Dual Degree (B.Tech. & M.Tech.) in Engineering Physics (Major) and Electrical Engg. (Minor).	(August 2015 – May 2021)    (July 2010 – August 2015)
AWARDS	Stanford SAIL Postdoctoral Fellowship, 2020 Simons Berkeley Research Fellowship (declined), 2020 Institute Silver Medal, IIT Bombay, 2015 Institute Academic Prize, IIT Bombay, 2012–2014	
WORK EXPERIENCE	<b>Research Intern, Google Brain</b> Advisors: Behnam Neyshabur & Hanie Sedghi	(Summer 2019)
PREPRINTS	<p>Understanding Scaling Laws: Sample and model size tradeoffs in two-layer neural networks. Niladri Chatterji, Percy Liang, Tatsunori Hashimoto. In preparation, 2022.</p> <p>Holistic Evaluation of Language Models. Percy Liang and Niladri Chatterji et al. arXiv preprint, 2022. Under review at TMLR.</p> <p>Deep linear networks can benignly overfit when shallow ones do. Niladri Chatterji, Philip Long. arXiv preprint, 2022. Under review at JMLR.</p>	

Undersampling is a minimax optimal robustness intervention in nonparametric classification.  
Niladri Chatterji\*, Saminul Haque\*, Tatsunori Hashimoto.  
arXiv preprint, 2022.

Random feature amplification: Feature learning and generalization in neural networks.  
Spencer Frei, Niladri Chatterji, Peter Bartlett.  
arXiv preprint, 2022. Under review at JMLR.

On the opportunities and risks of foundation models.  
Rishi Bommasani and Niladri Chatterji et al.  
arXiv preprint, 2022. Under review at JMLR.

Convergence rates for Langevin Monte Carlo in the nonconvex setting.  
Xiang Cheng, Niladri Chatterji, Yasin Abbasi Yadkori, Peter Bartlett, Michael Jordan.  
arXiv preprint, 2018. Under review at JMLR.

## JOURNAL PAPERS

The interplay between implicit bias and benign overfitting in two-layer linear networks.  
Niladri Chatterji, Philip Long, Peter Bartlett.  
Journal of Machine Learning Research (JMLR), 2022.

Foolish crowds support benign overfitting.  
Niladri Chatterji, Philip Long.  
Journal of Machine Learning Research (JMLR), 2022.

Oracle lower bounds for sampling algorithms.  
Niladri Chatterji, Peter Bartlett, Philip Long.  
Bernoulli, 2022.

Is there an analog of Nesterov acceleration for gradient-based MCMC?  
Yi-An Ma, Niladri Chatterji, Xiang Cheng, Nicolas Flammarion, Peter Bartlett, Michael Jordan.  
Bernoulli, 2021.

When does gradient descent with logistic loss find interpolating two-layer networks?  
Niladri Chatterji, Philip Long, Peter Bartlett.  
Journal of Machine Learning Research (JMLR), 2021.

Finite-sample analysis of interpolating linear classifiers in the overparameterized regime.  
Niladri Chatterji, Philip Long.  
Journal of Machine Learning Research (JMLR), 2021.

Enhancement of spin-transfer torque switching via resonant tunneling  
Niladri Chatterji, Ashwin Tulapurkar, Bhaskaran Muralidharan.  
Applied Physics Letters, 2014.

## CONFERENCE PAPERS

Benign overfitting without linearity: Neural network classifiers trained by gradient descent for noisy linear data.  
Spencer Frei, Niladri Chatterji, Peter Bartlett.  
Conference on Learning Theory (COLT), 2022.

Is importance weighting incompatible with interpolating classifiers?  
Ke Alexander Wang\*, Niladri Chatterji\*, Saminul Haque, Tatsunori Hashimoto.  
International Conference on Learning Representations (ICLR), 2022.

On the theory of reinforcement learning with once-per-episode feedback.  
Niladri Chatterji\*, Aldo Pacchiano\*, Peter Bartlett, Michael Jordan.  
Advances in Neural Information Processing Systems (NeurIPS), 2021.

When does gradient descent with logistic loss interpolate using deep networks with smoothed ReLU activations?  
Niladri Chatterji, Philip Long, Peter Bartlett.  
Conference on Learning Theory (COLT), 2021.

OSOM: A simultaneously optimal algorithm for multi-armed and linear contextual bandits.  
Niladri Chatterji, Vidya Muthukumar, Peter Bartlett.  
International Conference on Artificial Intelligence and Statistics (AISTATS), 2020.

Langevin Monte Carlo without smoothness.  
Niladri Chatterji\*, Jelena Diakonikolas\*, Michael Jordan, Peter Bartlett.  
International Conference on Artificial Intelligence and Statistics (AISTATS), 2020.

The intriguing role of module criticality in the generalization of deep networks.  
Niladri Chatterji, Behnam Neyshabur, Hanie Sedghi.  
International Conference on Learning Representations (ICLR), 2020. **(Spotlight Talk)**

Online learning with kernel losses.  
Niladri Chatterji\*, Aldo Pacchiano\*, Peter Bartlett.  
International Conference on Machine Learning (ICML), 2019. **(Long Talk)**

Underdamped Langevin MCMC: A non-asymptotic analysis.  
Xiang Cheng\*, Niladri Chatterji\*, Peter Bartlett, Michael Jordan.  
Conference on Learning Theory (COLT), 2018.

On the theory of variance reduction for stochastic gradient Monte Carlo.  
Niladri Chatterji, Nicolas Flammarion, Yi-An Ma, Peter Bartlett, Michael Jordan.  
International Conference on Machine Learning (ICML), 2018.

Alternating minimization for dictionary learning: Local convergence guarantees.  
Niladri Chatterji, Peter Bartlett.  
Advances in Neural Information Processing Systems (NeurIPS), 2017.

## THESIS

Why do gradient methods work in optimization and sampling?  
Niladri Chatterji.  
Ph.D. Thesis, University of California Berkeley, 2021.

## SERVICE

### Journal Review

Journal of Machine Learning Research (JMLR), Bernoulli, Annals of Applied Probability, Annales de l'Institut Henri Poincaré (B) Probabilités et Statistiques, SIAM Journal on the Mathematics of Data Science.

### Conference Review

NeurIPS (2018–), ICML (2020–), COLT (2018–), AISTATS (2019, 2020), ALT (2018, 2021), ICLR (2021), L4DC (2020, 2021).

February 9, 2023